

```
In [1]: # import libraries
import pandas as pd
import numpy as np
```

```
In [76]: # Load data
df= pd.read_csv('hr_dataset.csv')
df.head()
```

Out[76]:

	Employee_ID	Age	Department	Satisfaction_Level	Last_Evaluation	Projects	Average_Mon
0	896999	NaN	Finance	0.42	0.08	NaN	
1	331148	NaN	Hr	0.91	0.73	NaN	
2	559437	36.0	Operations	0.93	0.82	7.0	
3	883201	41.0	Finance	0.03	0.53	7.0	
4	562242	NaN	Finance	0.66	0.72	3.0	

```
In [77]: print(df.head(40))
```

## Notebooks

	Employee_ID	Age	Department	Satisfaction_Level	Last_Evaluation	\
0	896999	NaN	Finance	0.42	0.08	
1	331148	NaN	Hr	0.91	0.73	
2	559437	36.0	Operations	0.93	0.82	
3	883201	41.0	Finance	0.03	0.53	
4	562242	NaN	Finance	0.66	0.72	
5	538510	NaN	Sales	0.85	0.76	
6	585585	36.0	NaN	0.65	0.39	
7	689574	NaN	Hr	0.89	0.80	
8	394433	58.0	Hr	0.33	0.92	
9	314638	NaN	IT	0.57	0.42	
10	9767	NaN	Hr	0.62	0.26	
11	747950	NaN	Finance	0.71	0.09	
12	102408	NaN	finanCe	0.27	0.12	
13	726713	NaN	NaN	0.31	0.00	
14	40517	43.0	NaN	0.13	0.68	
15	420135	NaN	NaN	0.65	0.21	
16	913501	NaN	Sales	0.28	0.80	
17	783720	NaN	Operations	0.44	0.88	
18	268678	NaN	NaN	0.09	0.56	
19	556615	NaN	Sales	0.13	0.02	
20	574096	NaN	Hr	0.09	0.75	
21	873303	48.0	Hr	0.95	0.70	
22	19284	23.0	finanCe	0.03	0.20	
23	323130	34.0	NaN	0.42	0.99	
24	189340	NaN	NaN	0.67	0.29	
25	776669	NaN	Hr	0.82	0.96	
26	634699	NaN	Hr	0.01	0.08	
27	110188	NaN	NaN	0.06	0.98	
28	196980	NaN	finanCe	0.17	0.60	
29	35597	NaN	Finance	0.76	0.05	
30	23450	NaN	Operations	0.20	0.29	
31	606051	31.0	NaN	0.12	0.00	
32	101643	33.0	Finance	0.29	0.08	
33	922195	57.0	Operations	0.20	0.53	
34	857149	44.0	finanCe	0.21	0.14	
35	827733	27.0	Hr	0.38	0.75	
36	857918	NaN	Finance	0.72	0.83	
37	68193	NaN	Hr	0.94	0.73	
38	94198	39.0	finanCe	0.01	0.69	
39	281086	NaN	Hr	0.42	0.96	

	Projects	Average_Monthly_Hours	Years_at_Company	Left
0	NaN	999	NaN	1.0
1	NaN	180	7.0	1.0
2	7.0	999	4.0	1.0
3	7.0	297	NaN	1.0
4	3.0	186	5.0	NaN
5	4.0	999	NaN	1.0
6	NaN	999	5.0	NaN
7	NaN	227	NaN	1.0
8	3.0	205	7.0	NaN
9	NaN	227	NaN	1.0
10	NaN	999	1.0	1.0
11	NaN	268	7.0	NaN
12	NaN	999	NaN	NaN
13	5.0	999	8.0	1.0
14	NaN	999	8.0	NaN
15	2.0	193	8.0	1.0
16	2.0	210	1.0	1.0

17	5.0	999	NaN	0.0
18	6.0	999	NaN	1.0
19	NaN	999	7.0	NaN
20	NaN	292	4.0	NaN
21	NaN	260	2.0	NaN
22	NaN	276	NaN	0.0
23	NaN	209	5.0	NaN
24	7.0	246	2.0	NaN
25	NaN	164	NaN	0.0
26	3.0	249	4.0	0.0
27	3.0	157	NaN	0.0
28	5.0	999	NaN	1.0
29	5.0	999	NaN	NaN
30	NaN	999	7.0	1.0
31	NaN	999	5.0	NaN
32	NaN	240	NaN	0.0
33	NaN	999	NaN	NaN
34	NaN	209	6.0	NaN
35	4.0	999	NaN	0.0
36	7.0	999	7.0	0.0
37	NaN	222	NaN	0.0
38	2.0	254	NaN	1.0
39	NaN	999	3.0	0.0

## DATA CLEANING

In [78]: `df.shape`

Out[78]: (30100, 9)

In [79]: `# check for null values  
df.isnull().sum()`

Out[79]: Employee\_ID 0  
Age 15075  
Department 3912  
Satisfaction\_Level 0  
Last\_Evaluation 0  
Projects 15052  
Average\_Monthly\_Hours 0  
Years\_at\_Company 15056  
Left 10005  
dtype: int64

In [81]: `# calculate the percentage of missing values for each column  
missing_counts = df.isnull().sum()  
missing_percent = (missing_counts / len(df)) * 100  
missing_percent = missing_percent.round(1)  
missing_percent`

```
Out[81]: Employee_ID      0.0
          Age            50.1
          Department     13.0
          Satisfaction_Level  0.0
          Last_Evaluation   0.0
          Projects         50.0
          Average_Monthly_Hours 0.0
          Years_at_Company 50.0
          Left             33.2
          dtype: float64
```

```
In [82]: # handle missing values
# Normalize text in column Departments, remove spaces and ensure the case is lower
df['Department']=df['Department'].str.lower().str.strip()
df['Department']= df['Department'].fillna('unknown')
```

```
In [83]: df.columns
```

```
Out[83]: Index(['Employee_ID', 'Age', 'Department', 'Satisfaction_Level',
       'Last_Evaluation', 'Projects', 'Average_Monthly_Hours',
       'Years_at_Company', 'Left'],
       dtype='object')
```

```
In [84]: # fill missing values in numerical columns with median
numerical_cols= ['Age', 'Projects', 'Years_at_Company', 'Left']
for col in numerical_cols:
    median_value = df[col].median()
    df[col]= df[col].fillna(median_value)
```

```
In [85]: df.isnull().sum()
```

```
Out[85]: Employee_ID      0
          Age            0
          Department     0
          Satisfaction_Level  0
          Last_Evaluation   0
          Projects         0
          Average_Monthly_Hours 0
          Years_at_Company 0
          Left             0
          dtype: int64
```

1. For Department, I chose to fill missing values with "unknown" rather than dropping rows.

This preserves all employees in the dataset (~13% had missing departments) and allows analysis of how employees with unknown departments compare to others.

Dropping rows would have removed valuable data and reduced the dataset unnecessarily.

2. For numerical columns (Age, Projects, YearsAtCompany, Left), I fill missing values with the median.

This is because the median is not affected by outliers, and it maintains the central tendency of the data without introducing bias from extreme values.

In [86]:

```
df.describe()
```

Out[86]:

	<b>Employee_ID</b>	<b>Age</b>	<b>Satisfaction_Level</b>	<b>Last_Evaluation</b>	<b>Projects</b>	<b>Average</b>
<b>count</b>	30100.000000	30100.000000	30100.000000	30100.000000	30100.000000	30100.000000
<b>mean</b>	501459.486844	41.021528	0.500229	0.500954	4.751595	
<b>std</b>	289871.790946	7.925980	0.287933	0.290158	1.232629	
<b>min</b>	1005.000000	22.000000	0.000000	0.000000	2.000000	
<b>25%</b>	250675.750000	41.000000	0.250000	0.250000	5.000000	
<b>50%</b>	500246.500000	41.000000	0.500000	0.500000	5.000000	
<b>75%</b>	754074.000000	41.000000	0.750000	0.760000	5.000000	
<b>max</b>	999999.000000	60.000000	1.000000	1.000000	7.000000	

- Average\_Monthly\_Hours has unrealistic values (e.g., 999 hours), which exceed the maximum possible monthly hours (~730). These could be placeholders for missing or invalid data and would distort analysis and model performance.

Fix: Cap values above 400 hours at 200 hours, a realistic monthly workload baseline ( $\approx 45$  hours/week  $\times$  4.3 weeks/month). This ensures the values within humanly possible limits without introducing artificial averages, preserving data integrity.

- Age (22–60 years): No outliers detected; all values fall within a realistic working-age range.
- Satisfaction\_Level & Last\_Evaluation (0.0–1.0): Values fall within expected normalized limits.
- Projects (2–7 projects): Range reflects a reasonable employee workload.
- Years\_at\_Company (1–10 years): Values fall within a typical tenure range.

In [87]:

```
df.loc[df['Average_Monthly_Hours'] > 300, 'Average_Monthly_Hours']=200
```

In [88]:

```
print("\nOutliers capped using human workload baseline. ")
```

Outliers capped using human workload baseline.

In [89]:

```
df.duplicated().value_counts()
```

```
Out[89]: False    30000
          True     100
          Name: count, dtype: int64
```

```
In [90]: #Remove Exact Duplicate Rows
df = df.drop_duplicates()
print("\nDuplicate rows removed .")
```

```
Duplicate rows removed .
```

```
In [91]: #Ensure Unique Employee_IDs
# Remove duplicate Employee_IDs (keep first valid record)
df = df.drop_duplicates(subset=["Employee_ID"], keep="first")
```

```
In [92]: df.duplicated().value_counts()
```

```
Out[92]: False    29524
          Name: count, dtype: int64
```

```
In [93]: df.duplicated().sum()
```

```
Out[93]: np.int64(0)
```

```
In [94]: df.duplicated(subset= ['Employee_ID']).sum()
```

```
Out[94]: np.int64(0)
```

```
In [95]: df.columns
```

```
Out[95]: Index(['Employee_ID', 'Age', 'Department', 'Satisfaction_Level',
                 'Last_Evaluation', 'Projects', 'Average_Monthly_Hours',
                 'Years_at_Company', 'Left'],
                 dtype='object')
```

```
In [96]: df= df.rename(columns={
          'Employee_ID': 'EmployeeID',
          'Age': 'Age',
          'Department': 'Department',
          'Satisfaction_Level': 'SatisfactionScore',
          'Last_Evaluation': 'LastEvaluationScore',
          'Projects': 'NumProjects',
          'Average_Monthly_Hours': 'AvgMonthlyHours',
          'Years_at_Company': 'YearsAtCompany',
          'Left': 'Attrition'})
```

```
In [97]: df.dtypes
```

```
Out[97]: EmployeeID      int64
          Age            float64
          Department     object
          SatisfactionScore   float64
          LastEvaluationScore   float64
          NumProjects      float64
          AvgMonthlyHours    int64
          YearsAtCompany    float64
          Attrition        float64
          dtype: object
```

After checking the data types:

- Convert 'Age', 'NumProjects', and 'YearsAtCompany' from float to integer because these values are whole numbers.
- Convert 'Attrition' to integer to clearly represent the binary target.
- Convert 'Department' to categorical to reflect discrete groups and improve memory efficiency.

```
In [98]: df["Age"] = df["Age"].astype("Int64")
          df["NumProjects"] = df["NumProjects"].astype("Int64")
          df["YearsAtCompany"] = df["YearsAtCompany"].astype("Int64")
          df["Attrition"] = df["Attrition"].astype("Int64")
```

```
In [99]: df["Department"] = df["Department"].astype("category")
```

```
In [101...]: print("\n---DATA QUALITY REPORT ---")
           print("Dataset shape:", df.shape)
           print("Employee_ID unique?", df["EmployeeID"].is_unique)
           print("\nRemaining missing values:\n", df.isnull().sum())
           print("\nDepartment distribution:\n", df["Department"].value_counts())
```

```
---DATA QUALITY REPORT ---
```

```
Dataset shape: (29524, 9)
```

```
Employee_ID unique? True
```

```
Remaining missing values:
```

```
EmployeeID      0  
Age            0  
Department     0  
SatisfactionScore 0  
LastEvaluationScore 0  
NumProjects    0  
AvgMonthlyHours 0  
YearsAtCompany 0  
Attrition      0  
dtype: int64
```

```
Department distribution:
```

```
Department  
finance      7314  
hr          7269  
unknown     3828  
sales        3790  
operations   3720  
it           3603  
Name: count, dtype: int64
```

```
In [103...]
```

```
## Save the cleaened data as csv  
df.to_csv("hr_cleaned_dataset.csv", index=False)  
print("\nCleaned HR dataset saved.")
```

```
Cleaned HR dataset saved.
```

```
In [ ]:
```